

How the JVM Executes Java

this slide intentionally left blank

Classes and Class Files

- Use javap to examine class file

HelloWorld0.java

```
public class HelloWorld0 {  
    public static void main ( String [] args ) {  
        System.out.println("Hello, it's Java  
time");  
    }  
}
```

```
$ javap HelloWorld0  
Compiled from "HelloWorld0.java"  
public class HelloWorld0 {  
    public HelloWorld0();  
    public static void main(java.lang.String[]);  
}
```

Classes and Class Files

- Use javap to examine class file

HelloWorld0.java

```
public class HelloWorld0 {
    public static void main ( String [] args ) {
        System.out.println("Hello, it's Java
time");
$ javap -c HelloWorld0
Compiled from "HelloWorld0.java"
public class HelloWorld0 {
    public HelloWorld0();
    Code:
        0: aload_0
        1: invokespecial #1    // Method java/lang/Object."<init>":()V
        4: return

    public static void main(java.lang.String[]);
    Code:
        0: getstatic     #2      // ...
        3: ldc          #3      // String Hello, it's Java time
        5: invokevirtual #4      // ...
        8: return
}
```

Classes and Class Files


- Use javap to examine class file

HelloWorld0.java

```
public class HelloWorld0 {
    public static void main ( String [] args ) {
        System.out.println("Hello, it's Java
time");
$ javap -verbose HelloWorld0
Classfile /Users/george/work/java/JVM/HelloWorld0.class
  Last modified 27-May-2014; size 437 bytes
  MD5 checksum d587c1c612c2809f3a6288317fd631fa
  Compiled from "HelloWorld0.java"
public class HelloWorld0
  SourceFile: "HelloWorld0.java"
  minor version: 0
  major version: 52
  flags: ACC_PUBLIC, ACC_SUPER
Constant pool:
   #1 = Methodref   #6.#15    // java/lang/Object."<init>":()V
   #2 = Fieldref    #16.#17   // java/lang/System.out:Ljava/io...
   #3 = String      #18      // Hello, it's Java time
   #4 = Methodref   #19.#20   // ...
   ...
```

Introducing Bytecode

- Executable code for the Java Virtual Machine
- Compact – 1 byte for opcode
 - most opcodes do not take arguments
 - JVM is stack based machine
- Opcodes are largely type related
 - fload float
 - imul int
 - dstore double
- Also control flow
 - invokevirtual
 - goto
 - areturn



Introducing Bytecode


- A simple example:

```
public class IncExample {  
    public int inc ( int i ) {  
        int result;  
        result = i + 1;  
        return result;  
    }  
}
```

```
$ javap -verbose IncExample  
...  
public int inc(int);  
  descriptor: (I)I  
  flags: ACC_PUBLIC  
  Code:  
    stack=2, locals=3, args_size=2  
    0: iload_1  
    1: iconst_1  
    2: iadd  
    3: istore_2  
    4: iload_2  
    5: ireturn  
  LineNumberTable:  
    line 4: 0  
    line 5: 4  
...
```

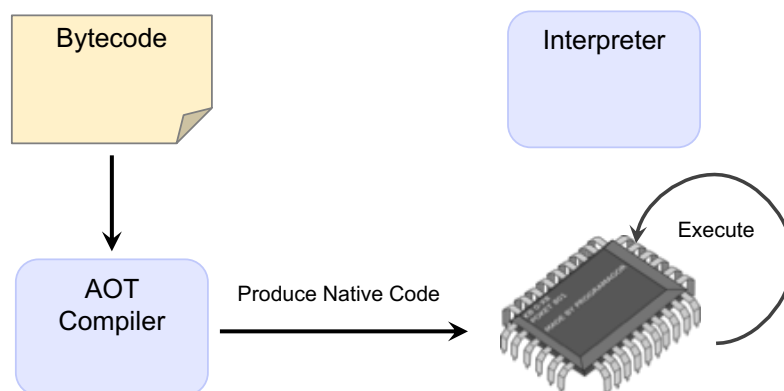
Executing Bytecode

- The simple view



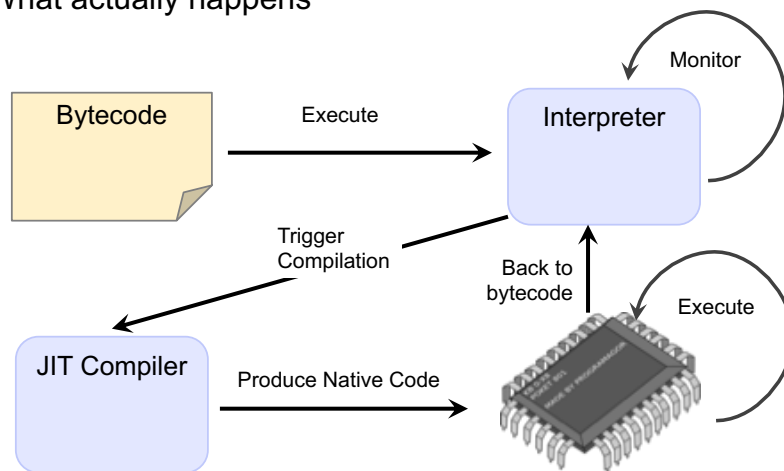
Executing Bytecode

- The slightly more complicated view



Executing Bytecode

- What actually happens



Dynamic Compilation Options

- Client mode
 - `java -client`
 - less aggressive optimisation
 - no speculative optimisation
 - smaller memory footprint
 - faster startup
- Server mode
 - `java -server`
 - full optimisation
 - suited to long running applications
- Tiered Compilation
 - start by using client compilations
 - switch to server mode when available

Execution Example

- Simple example
 - loop necessary to see dynamic compilation

```
public class HelloWorld {  
  
    public static void main(String[] args) {  
        for (int i = 0; i < 100_000; i++) {  
            hello();  
        }  
    }  
  
    private static void hello() {  
        System.err.println("Hello, world!");  
    }  
}
```

Execution Example

- Run in default mode

```
$ time java HelloWorld 2>/dev/null  
real    0m0.336s  
user    0m0.467s  
sys     0m0.120s
```

- Interpreted mode

```
$ time java -Xint HelloWorld 2> /dev/null  
real    0m1.969s  
user    0m1.828s  
sys     0m0.152s
```

- Compiled mode

```
$ time java -Xcomp HelloWorld 2> /dev/null  
real    0m2.365s  
user    0m2.198s  
sys     0m0.241s
```

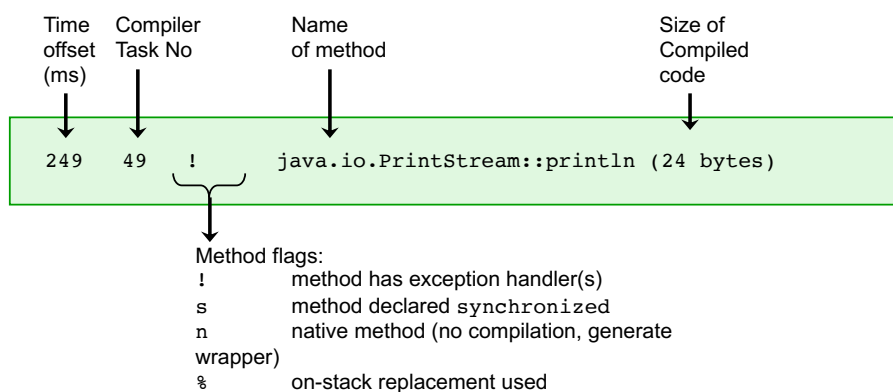
Dynamic Compiler Diagnostics

- Command line flags provide further information

```
$ java -XX:-TieredCompilation -XX:+PrintCompilation HelloWorld
2> /dev/null
72    1      java.lang.String::hashCode (55 bytes)
79    2      java.lang.String::indexOf (70 bytes)
104   3      sun.nio.cs.UTF_8$Encoder::encodeArrayLoop
                                (489 bytes)
120   4      java.nio.Buffer::position (5 bytes)
122   5      java.nio.ByteBuffer::arrayOffset (35 bytes)
127   6      java.nio.Buffer::position (43 bytes)
127   7      n  java.lang.System::arraycopy (native) (static)
...
247  47      java.lang.String::indexOf (7 bytes)
247  48      HelloWorld::hello (9 bytes)
249  49      !  java.io.PrintStream::println (24 bytes)
256  50      java.io.PrintStream::print (13 bytes)
259  51      !  java.io.PrintStream::write (83 bytes)
262  52      !  java.io.PrintStream::newLine (73 bytes)
264  53      java.io.BufferedWriter::newLine (9 bytes)
265  54      %  HelloWorld::main @ 2 (18 bytes)
```

Dynamic Compiler Diagnostics

- Decoding the output



Dynamic Compiler Diagnostics

- `-Xbatch` forces compiler into application thread
 - application paused while compiler runs
 - `b` flag shown for compilation
 - overall time increases

```
$ java -Xbatch -XX:-TieredCompilation -XX:+PrintCompilation
HelloWorld 2> /dev/null

75   1   b   java.lang.String::hashCode (55 bytes)
84   2   b   java.lang.String::indexOf (70 bytes)
112  3   b   sun.nio.cs.UTF_8$Encoder::encodeArrayLoop
                                     (489 bytes)
141  4   b   java.nio.Buffer::position (43 bytes)
143  5   n   java.lang.System::arraycopy (native)   (static)
147  6   b   java.nio.Buffer::position (5 bytes)
161  7   b   java.nio.charset.CoderResult::isUnderflow (13 bytes)
161  8   b   java.io.BufferedWriter::ensureOpen (18 bytes)
162  9   b   java.io.PrintStream::ensureOpen (18 bytes)
162 10  !b   java.io.BufferedWriter::write (117 bytes)
...
```

Dynamic Compiler Diagnostics

- Summary statistics can be obtained

```
$ java -Xbatch -XX:-TieredCompilation -XX:+PrintCompilation
-XXX:+CITime HelloWorld 2> /dev/null
...
Accumulated compiler times (for compiled methods only)
-----
Total compilation time      : 0.195 s
Standard compilation       : 0.194 s, Average : 0.004
On stack replacement       : 0.001 s, Average : 0.001

Total compiled methods      : 52 methods
Standard compilation       : 51 methods
On stack replacement       : 1 methods
Total compiled bytecodes    : 11838 bytes
Standard compilation       : 11820 bytes
On stack replacement       : 18 bytes
Average compilation speed   : 60754 bytes/s

nmethod code size          : 57376 bytes
nmethod total size         : 111752 bytes
```


On Stack Replacement?

- Normal compilation triggered by call count
 - method has been called a specific number of times
 - `-XX:CompileThreshold`
 - defaults to 10000
- JVM can detect method that is in long loop
 - e.g. `main` in example
- Compiler can replace method while it is executing
 - On Stack Replacement

```
422    54  % b      HelloWorld::main @ 2 (18 bytes)
```

LogCompilation

- For more detail than `PrintCompilation` provides
- Generates a detailed (& large) XML file
 - Often 100s of MB
- Full detail of compilation events
- Hard to handle without tooling (e.g. `JITWatch`)

Viewing Assembly Code

- Possible to see assembly code output from compiler
 - requires additional library for display

```
public class TinyExample {  
  
    public static void main ( String [] args ) {  
        for ( int i=0; i < 1000000; i++ ) {  
            tiny();  
        }  
    }  
  
    private static int tiny() {  
        return 1 + 1;  
    }  
}
```

Assembly Primer

- AT&T format
 - Usually <operator> <src>, <dst>

```
mov %rdx, %rax    ; move %rdx into accumulator  
add %rcx, %rax    ; add %rcx to accumulator  
XOR %eax, %eax    ; Zero a register  
mov $0, %eax      ; Zero a register  
XCHG %eax, %eax   ; No-op, often used for memory fences
```

Assembly Syntax

- Common prefixes

b – byte (8 bits)
s – short (16 bits)
w – word (16 bits)
l – long (32 bits)
q – quad (64 bits)
t – ten (80 bits, floating point)

ax – accumulator
bp – frame pointer
sp – stack pointer

e – 32-bit
r – 64-bit

Viewing Assembly Code

```
$ java -XX:+UnlockDiagnosticVMOptions -XX:+PrintAssembly TinyExample
...
Code:
[Entry Point]
[Verified Entry Point]
[Constants]
# {method} {0x0000000123b9f300} 'tiny' '()'I in 'TinyExample'
# [sp+0x40] (sp of caller)
0x000000010f211560: mov    %eax,-0x14000(%rsp)
0x000000010f211567: push   %rbp
0x000000010f211568: sub    $0x30,%rsp          ;*iconst_2
                                ; -
TinyExample::tiny@0 (line 10)

0x000000010f21156c: mov    $0x2,%eax
0x000000010f211571: add    $0x30,%rsp
0x000000010f211575: pop    %rbp
0x000000010f211576: test   %eax,-0x202d47c(%rip) #
0x000000010d1e4100                                ; {poll_return}

0x000000010f21157c: retq

...
```

Basic Optimisation

- Method inlining
 - most commonly applied optimisation
 - removes call overhead
 - can be tuned with command line options

```
$ java -Xbatch -XX:-TieredCompilation -XX:+PrintCompilation
-XX:+UnlockDiagnosticVMOptions -XX:+PrintInlining
HelloWorld 2> /dev/null

...
191 14  b  java.io.BufferedOutputStream::flushBuffer (29 bytes)
        @ 20  java.io.FileOutputStream::write (12 bytes)
                        inline (hot)
        \-> TypeProfile (4466/4466 counts) =
                        java/io/FileOutputStream
        @ 8   java.io.FileOutputStream::writeBytes (0 bytes)
                        native method
...
```


Monomorphic Dispatch

- How many different types are seen at a call site?
 - often, it's only 1
 - JVM optimizes for this case
 - aggressive optimisation (so only in server mode)
 - can be backed out
- Hotspot optimizes vtbl lookup
 - subclasses have the same vtbl structure as their parent
 - hotspot collapses the child vtbl into the parent
 - class word in the object header is checked
 - if changed then this optimisation is backed out
- Classloading can invalidate monomorphic dispatch

Further JVM Optimisations

- Loop Unrolling

```
private static final String[] options =  
    { "Yes", "No", "Cancel" };  
  
public void meth1 () {  
    for ( String opt: options ) {  
        process(opt);  
    }  
}
```




```
private static final String[] options =  
    { "Yes", "No", "Cancel" };  
  
public void meth1 () {  
    process(options[0]);  
    process(options[1]);  
    process(options[2]);  
}
```

Further JVM Optimisations

- Lock Coarsening
 - effective but depends on scope of lock

```
public void doStuff () {  
    for ( String opt: options ) {  
        process(opt);  
    }  
}  
  
public synchronized String process ( String opt ) {  
    ...  
}
```



```
public void doStuff () {  
    synchronized ( this ) {  
        for ( String opt: options ) {  
            process(opt);  
        }  
    }  
}
```


Further JVM Optimisations

- Lock Elision

– also depends on local scope of lock

```
public void doStuff () {  
    List l = new ArrayList();  
    synchronized (l) {  
        for (option : options) {  
            l.add(process(option));  
        }  
    }  
}
```

List l is local
and does not
"escape"
this thread



```
public void doStuff () {  
    List l = new ArrayList();  
    for (option : options) {  
        l.add(process(option));  
    }  
}
```

Lock is
not
required

Further JVM Optimisations

- Dead Code Elimination

– remove code that does nothing

```
public class TinyExample {  
  
    public static void main ( String [] args ) {  
        for ( int i=0; i < 1000000; i++ ) {  
            tiny();  
        }  
    }  
  
    private static void tiny() {  
        int i = 100;  
        i += 1;  
    }  
}
```

Code has no effect,
may be removed
during compilation

Further JVM Optimisations

- Escape Analysis
 - allows reduction of heap usage

```
private static class Foo {  
    public final String a;  
    public final String b;  
  
    Foo(String a, String b) {  
        this.a = a;  
        this.b = b;  
    }  
}
```

```
public void one() {  
    Foo f = new Foo("Hello", "JVM");  
    two(f);  
}  
  
public void two(Foo f) {  
    System.out.print(f.a);  
    System.out.print(", ");  
    three(f);  
}  
  
public void three(Foo f) {  
    System.out.print(f.b);  
    System.out.println('!');  
}
```

Further JVM Optimisations

- Escape Analysis
 - allows reduction of heap usage

```
private static class Foo {  
    public final String a;  
    public final String b;  
  
    Foo(String a, String b) {  
        this.a = a;  
        this.b = b;  
    }  
}
```

```
public void oneTwoThree() {  
    System.out.print("Hello");  
    System.out.print(", ");  
    System.out.print("JavaOne");  
    System.out.println('!');  
}
```